

# Data Manipulation with Pandas

Hrachya Asatryan

# What This Lecture Is (and Isn't)

## **This lecture is about**

- Understanding data as an object
- How pandas represents and manipulates data
- How to *reason* about datasets

## **This lecture is NOT**

- A pandas API reference
- A tutorial on every method
- Machine learning

# Why Pandas Exists

Before pandas:

- NumPy → fast but structureless
- SQL → structured but rigid
- Excel → flexible but dangerous

Pandas exists to:

- Work with **structured, labeled data**
- Iterate quickly
- Preserve meaning

**Pandas is a library for labeled, tabular data and the operations that make sense on it.**

# What “Labeled” Means

- Rows and columns have **names**
- Operations align by **labels**, not positions
- Labels are part of the data

This changes how computations behave.

# Pandas Core Abstractions

# Two Fundamental Objects

Pandas has two core data structures:

- **Series** (1D, labeled)
- **DataFrame** (2D, labeled)

Everything else builds on these.

# pd.Series: Labeled 1D Data

A Series is:

- A 1D array
- With an **index**

```
s = pd.Series([72, 85, 90], index=["math", "physics", "cs"])
```

pd.Series is both an array and a mapping

- Array-like: supports vectorized operations
- Dictionary-like: keys → values

```
s["physics"]  
s.iloc[1]
```

**Index ≠ position.**

# DataFrame: Aligned Series

A DataFrame is:

- A collection of Series
- Sharing the same index

```
df["math"] # Series
```

Each column is its own labeled object.

# DataFrame

DataFrame named states\_df

integer position	label index	0	1	2	integer position
		text	capital	population	
0	'OH'	'Ohio'	Columbus'	11799448	
1	'TN'	'Tennessee'	'Nashville'	6910840	
2	'AZ'	'Arizona'	'Phoenix'	7151502	states_df.loc['AZ'] states_df.iloc[2]
3	'PA'	'Pennsylvania'	'Harrisburg'	13002700	series
4	'AK'	'Alaska'	'Juneau'	733391	states_df['capital'] states_df.capital

single cell value

- **Data frames** are essentially tables.
- The **values** of columns or rows are **series**.

# Columns vs Rows

Which axis do you think it is natural to do operations on?

# Columns vs Rows

Which axis do you think it is natural to do operations on?

- If I want to calculate statistics for the data, it makes sense to do it per-column
- Column-wise operations are natural
- Row-wise operations require intent (`axis`)

```
df.mean()
```

```
df.mean(axis=1)
```

This is not accidental.

# Indexing and Alignment

# The Index Is Not Optional

The index:

- Identifies data
- Controls alignment
- Can be anything (strings, dates, categories)

# Two Kinds of Indexing

- `.loc` → label-based
- `.iloc` → position-based

```
df.loc["Alice"]
```

```
df.iloc[0]
```

**Never confuse them.**

# Why Pandas Forces This Separation

- Labels and positions are not the same thing
- **Silent mistakes are worse than errors**
- Alignment is the core idea
  - Pandas aligns data by labels
- $s1 + s2$ 
  - Matching labels combine
  - Non-matching labels -> NaN

# Alignment Example

```
s1 = pd.Series([1,2,3], index=["a","b","c"])
s2 = pd.Series([10,20,30], index=["b","c","d"])
```

Result:

- Correct
- Dangerous
- Silent

# Alignment in DataFrames

```
df - df.mean()
```

- Mean is computed per column
- Subtraction aligns on column names

This is **intentional broadcasting**.

# Pandas Will Not Save You

**Pandas assumes you know what you're doing.**

- No warnings
- No complaints
- Just NaNs

# Data Types

## Common pandas dtypes

- int64, float64
- object
- category
- datetime64

If you see object, be suspicious. It usually means

“ pandas gave up on understanding your data ”

# Missing Data Is Information

Missing data:

- Is not zero
- Is not false
- Has meaning

Represented as **NaN**.

# Missing Data Operations

```
df.isna()  
df.dropna()  
df.fillna(0)
```

**Handling missing data is a decision, not a fix.**

# Working with Data

# What is data?

## **Data Is Not Reality**

Data is:

- A measurement
- A proxy
- Often incomplete or biased

**Reality → measurement → data**

# Data Is a Table of Claims

**Every row is a claim about the world.**

**Every column is a type of claim.**

Example:

- Row: “This person exists”
- Column: “This person’s age is 23”

If a value is missing or wrong → the claim is weak or false.

# Data Isn't Just Numbers

Data includes:

- Measurements
- Categories
- Identifiers
- Timestamps
- Flags

Pandas treats these differently **on purpose**.

This is why `dtype` matters.

# Columns Are Not Equal

Columns play different roles:

- Identifiers (IDs)
- Features (measurements)
- Targets (labels)
- Metadata (timestamps, source)

Pandas doesn't know this – **you must**.

# Pandas Assumes Tabular Semantics

Pandas assumes:

- Rows are observations
- Columns are variables
- Columns have meaning across rows

This is why:

- vectorization works
- groupby makes sense
- alignment exists

# Pandas Assumes Comparability

If values are in the same column:

- They are comparable
- They share units
- They can be summarized

If this is false, **your analysis is invalid.**

# Pandas Is Column-Oriented

- Columns are first-class
- Operations default to column-wise
- Rows are secondary

This is why `df.mean()` works without arguments.

# Tall vs Wide Data

The diagram illustrates the difference between tall and wide data formats. The 'TALL' section on the left shows a hierarchical table where rows are grouped by state and loan type. The 'WIDE' section on the right shows a flat table where each row represents a unique combination of state and loan type.

State Name	Loan Type Name	Avg Conventional Loan Amount	Avg FHA-Insured Loan Amount	Avg FSA/RHS-guaranteed Loan Amount	Avg VA-guaranteed Loan Amount
Null	Conventional	187,774			
	FHA-insured	141,054			
	FSA/RHS-guaranteed	109,502			
	VA-guaranteed	192,631			
Alabama	Conventional	195,334			
	FHA-insured	134,300			
	FSA/RHS-guaranteed	121,052			
	VA-guaranteed	197,487			
Alaska	Conventional	264,511			
	FHA-insured	245,730			
	FSA/RHS-guaranteed	201,451			
	VA-guaranteed	302,465			
Arizona	Conventional	246,171			
	FHA-insured	172,454			
	FSA/RHS-guaranteed	149,408			
	VA-guaranteed	226,221			
Arkansas	Conventional	173,943			
	FHA-insured	129,261			
	FSA/RHS-guaranteed	110,871			
	VA-guaranteed	169,422			
California	Conventional	471,505			
	FHA-insured	314,254			
	FSA/RHS-guaranteed	241,561			
	VA-guaranteed	348,666			

Pandas can handle both, but GroupBy prefers tall data while Visualization prefers tall data.

# One Variable per Column

Rule:

**One column = one variable.**

Violations:

- “Jan\_sales”, “Feb\_sales”, ...
- Multiple units in one column
- Encoded categories in strings

These force awkward pandas code later.

# Index as a Coordinate System

Index defines:

- What makes a row unique
- How data aligns
- What joins mean

Index is **not decoration**.

# The Real Data Loop

**Load → Inspect → Question → Transform → Summarize → Repeat**

not

**Load → Model → Profit**

# Inspect Before You Touch Anything

Mandatory first steps:

`df.shape`

`df.columns`

`df.dtypes`

`df.head()`

`df.info()`

If you skip this, everything after is suspect.

# Questions Come First, Code Second

Bad workflow:

**“Which pandas method should I use?”**

Good workflow:

**“What question am I answering?”**

Pandas methods are answers, not goals.

# Typical Questions and Pandas Thinking

# Question: “Are Groups Different?”

Translation:

- Identify grouping variable
- Summarize numeric columns

```
df.groupby("group").mean()
```

# Question: “Is Something Unusual?”

Translation:

- Look at ranges
- Look at counts
- Look at missingness

```
df.describe()
```

# Question: “Does This Change Over Time?”

Translation:

- Time-based index
- Sorting
- Aggregation

Pandas is designed for this.

# Question: “Do Variables Move Together?”

Translation:

- Numeric columns
- Summaries or correlations

Later → visualization.

# Why Visualization Comes Next

Tables hide:

- Distribution
- Outliers
- Structure

Visualization reveals them.

# Summary

- Pandas is about **labeled data**
- Indexing and alignment matter
- Data analysis is iterative
- Pandas helps you ask questions

# What Comes Next

- Exploratory Data Analysis
- Visualization
- Seeing patterns instead of guessing

# Questions?

# Jupyter Notebook Demo

In [this jupyter notebook](#), we will be demoing the core ideas of Pandas: indexing and alignment, as well as inspecting and questioning data.

# Resources

- [Pandas Official Documentation](#)
- [Jake VanderPlas Python Data Science Handbook](#)